

Uno: a composable framework for nonlinearly constrained optimization

Charlie Vanaret^{1¶} and Alexis Montoisson^{2*}

¹ Applied Optimization Department, Zuse-Institut Berlin, Germany ² Mathematics and Computer Science Division, Argonne National Laboratory, USA ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Uno is a composable software framework for nonlinearly constrained optimization written in modern C++. It unifies the workflows of Lagrange-Newton methods, i.e., gradient-based algorithms that iteratively solve the KKT optimality conditions using Newton's method. The central idea is to decompose these methods into reusable, interchangeable components (constraint reformulation, step computation, globalization techniques, and acceptance criteria) so that classical and hybrid algorithms can be assembled, compared, and tested within a single framework rather than reimplemented as separate solvers. As of June 2026, Uno supports sequential (convex and nonconvex) quadratic programming, interior-point (barrier) methods, sequential linear programming, and unconstrained optimization. For full mathematical details of the algorithms implemented in Uno, see ([Vanaret & Leyffer, 2026](#)).

Uno has interfaces to Julia, Python, C, Fortran, and AMPL, enabling interoperability across scientific computing environments. Precompiled artifacts are available on GitHub, and the solver can be accessed directly via `UnoSolver.jl` in Julia or `unopy` in Python.

Statement of need

Nonlinearly constrained optimization is central to engineering, optimal control ([Lewis et al., 2012](#)), machine learning ([Sra et al., 2011](#)), scientific modeling ([Nocedal & Wright, 2006](#)), and plays a key role in mixed-integer nonlinear optimization ([Lee & Leyffer, 2011](#)). The major solution paradigms (sequential quadratic programming, interior-point methods, and augmented Lagrangian methods) are usually developed and implemented as separate solver families. Yet they share the same building blocks: step computation, constraint handling, Hessian models, and globalization strategies. Because these building blocks are rigid inside monolithic codes, algorithmic ideas are typically tested at the level of *complete solvers* rather than *building blocks*: evaluating a new globalization strategy or Hessian approximation means reimplementing an entire solver or intrusively modifying an existing one.

Uno targets two audiences. For **algorithm developers and optimization researchers**, it makes the building blocks explicit and recombina- ble across paradigms, so that a single new strategy can be swapped into an otherwise state-of-the-art method and benchmarked in isolation. For **practitioners and end users**, it exposes those building blocks, pre-assembled into presets that reproduce established solvers, behind multiple language interfaces. Because these presets live in a single framework with a common interface, users can identify the most suitable method for a given class of problems by reconfiguring strategies within Uno, rather than installing and benchmarking several independent solvers.

*Now at Advanced Micro Devices (AMD).

39 The gap Uno fills is the absence of a framework in which *production-grade* algorithmic
40 components can be composed across solver paradigms under a unified abstraction.

41 State of the field

42 Popular solvers such as IPOPT (Wächter & Biegler, 2006), KNITRO (Byrd et al., 2006),
43 SNOPT (Gill et al., 2005), and WORHP (Büskens & Wassel, 2012) are robust and efficient
44 but monolithic: parameter tuning is exposed while internal components remain rigid. They
45 are designed for end users rather than algorithmic experimentation, and testing a new variant
46 usually requires intrusive modification or reimplementation. A notable exception among
47 production codes is MadNLP.jl (Shin et al., 2024), an actively developed Julia solver, though
48 it remains focused on interior-point methods.

49 Frameworks that do emphasize modularity occupy a different niche. General-purpose libraries
50 such as `scipy.optimize`, `NLopt`, `pyOpt`, and `pyOptSparse` let users select among complete
51 solvers, but they do not expose the *internal* algorithmic components for recombination. Closest
52 in spirit is `modOpt` (Joshy & Hwang, 2026), a Python environment that builds optimization
53 algorithms from self-contained modules (line searches, Hessian approximations, merit functions)
54 with an educational focus.

55 Uno is distinct on three counts. First, its components are production-grade algorithmic building
56 blocks compiled in modern C++. Second, it expresses these paradigms as instances of one
57 generic Lagrange-Newton method, so the same components are reused across all of them
58 rather than confined to a single method. Third, its presets reproduce state-of-the-art solvers at
59 competitive performance (see below), so a component swapped into Uno is measured against
60 a credible baseline rather than a teaching reference.

61 Software design

62 Within Uno, strategies are organized into a coherent hierarchy, illustrated in the wheel of
63 strategies (Figure 1). The outer ring lists the high-level *layers* of a Lagrange-Newton method
64 (for example, constraint reformulation and globalization); the middle ring lists the algorithmic
65 *ingredients* that Uno combines automatically to realize each layer; and the inner ring lists the
66 concrete *strategies* available for each ingredient. Strategies currently implemented in Uno are
67 highlighted in green.

68 In Uno's object-oriented architecture, the ingredients of Figure 1 are abstract classes whose
69 interfaces are implemented by subclasses (the strategies). Uno's simplified UML diagram is
70 shown in Figure 2: inheritance is represented as dashed lines with white arrows, composition
71 as solid lines with black diamonds; abstract classes are written in *italic* and subclasses in **bold**.
72 This architecture separates the mathematical logic of the algorithm (problem reformulation,
73 subproblem definition, globalization) from the computational aspects (evaluating the model's
74 functions, solving the subproblems).

75 Uno implements a generic Lagrange-Newton method in which the abstract classes interact and
76 exchange data while remaining agnostic to the underlying strategies. Strategies are selected
77 by the user at runtime via options. Uno also provides *presets*: particular combinations of
78 strategies and hyperparameters that correspond to state-of-the-art solvers. Two presets are
79 currently available: an `ipopt` preset mimicking the IPOPT solver (Wächter & Biegler, 2006)
80 (a *line-search restoration filter interior-point method with exact Hessian and primal-dual inertia*
81 *correction*), and a `filtersqp` preset mimicking the filterSQP solver (Fletcher & Leyffer, 1998)
82 (a *trust-region restoration filter SQP method with exact Hessian and no inertia correction*). In
83 principle all combinations of strategies can be generated, although some are not yet supported
84 (e.g., an interior-point method with a trust-region globalization mechanism). When no preset is
85 specified, Uno automatically dispatches to the `filtersqp` or `ipopt` preset based on the problem

size, sparing users the choice of the underlying method. Acting as an oracle, this selection mechanism can be readily extended to dispatch to new presets as they become available.




Figure 1: Unification framework: wheel of strategies.

These two presets perform on a par with the state-of-the-art solvers IPOPT (Uno is slightly less robust) and filterSQP (Uno is slightly more robust) in terms of function evaluations on a set of 429 small CUTE instances with fewer than 100 variables and constraints (Bongartz et al., 1995) translated to AMPL by Vanderbei. An up-to-date performance profile is maintained on Uno's documentation page. This benchmark shows that Uno's composability does not come at the cost of performance: assembling a method from modular components reproduces the behavior of the hand-written solvers it mimics.

Subproblem solvers are treated as interchangeable components that can be plugged in or swapped out without modifying the algorithmic logic, letting users match the solver to their problem structure, licensing constraints, or performance requirements. Uno currently interfaces several established LP, QP, and linear solvers: BQPD (Fletcher, 2000), HiGHS (Huangfu & Hall, 2018), MUMPS (Amestoy et al., 2000), MA27 (I. S. Duff & Reid, 1983), MA57 (Iain S. Duff, 2004), MA86 (J. Hogg & Scott, 2010), SSIDS (J. D. Hogg et al., 2016), and Krylov.jl (Montoison & Orban, 2023).

Interfaces

To make Uno accessible to a wide range of users, we provide multiple language interfaces.

The AMPL Solver Library (Gay, 1997) interface gives access to the AMPL (Fourer et al., 1989) modeling language for optimization. It is distributed as a binary that takes a compiled AMPL model (.nl file) as input, allowing users to solve problems without directly interacting with the C++ core.

The C interface provides direct access to Uno's core functionality while maximizing

interoperability with other programming languages and tools. The optimization process is expressed as the interaction between two independent opaque objects: a model describing the optimization problem and a solver holding the algorithmic configuration and execution state. The interface is therefore centered around two main structures:

- **Model:** represents an optimization problem and stores information about variables, bounds, constraints, the objective function, and derivative information. Users create a model and set its components (objective, constraints, derivatives, initial primal-dual point).
- **Solver:** represents the algorithm used to solve a given model. Users set solver options, attach callbacks, and access results such as primal and dual solutions, residuals, iteration counts, and performance metrics.

The Fortran interface provides access to Uno through the C API using `iso_c_binding`. It is split into two files: `uno_c.f90` for low-level C bindings, and `uno_fortran.f90` for Fortran-friendly wrappers that handle string arguments. The interface can be included directly in a source file or wrapped in a module for cleaner use statements. It is provided as source rather than a precompiled Fortran module (`.mod`) to maximize portability and interoperability across compilers and platforms.

The Julia interface is distributed as the registered package `UnoSolver.jl`. It integrates Uno into the Julia optimization ecosystem through:

- a thin wrapper around the full C API,
- an interface to `NLPModels.jl` for solving problems following the `NLPModels` API, such as `ADNLPModels.jl` or `ExaModels.jl`,
- an interface to `MathOptInterface.jl` for handling JuMP models.

Precompiled artifacts are downloaded automatically, making the package plug-and-play without requiring user compilation.

The Python interface is registered on PyPI as `unopy`, providing Uno bindings via precompiled wheels on most platforms.

MATLAB and Rust interfaces are also under development, further expanding Uno's accessibility.

Looking ahead, we plan to integrate Uno as a continuous relaxation solver within MINLP frameworks such as SCIP (Bestuzheva et al., 2023), which will require efficient reoptimization that reuses internal solver state (active sets, factorizations, preallocated workspace) across structurally similar solves.

Research impact statement

Uno has demonstrated value both as a research tool and as a dependency of established software. Its original authors used it to prototype a novel globalization strategy, the funnel method (Kiessling et al., 2025). It has since been used as the lower-level solver in derivative-free bilevel optimization (Cesaroni et al., 2026), been benchmarked on Euclidean problems with bound constraints (Baran et al., 2026), and been cited in (Cederberg et al., 2026; Desef, 2025; Gruss, 2024).

Uno is currently available as a nonlinear optimization solver in:

- Julia:
 - the JuMP.jl ecosystem (Dunning et al., 2017),
 - `control-toolbox`, a collection of Julia packages for mathematical control and its applications (Caillau et al., 2024),
- C++:
 - `CasADi`, an open-source tool for nonlinear optimization and algorithmic differentiation (Andersson et al., 2019),

- Python:
 - [pyOptSparse](#), an object-oriented framework for formulating and solving nonlinear constrained optimization problems ([Wu et al., 2020](#)),
 - [DNLP](#), an extension of CVXPY to general nonlinear programming ([Cederberg et al., 2026](#)),
- Fortran:
 - [CUTEst](#), the Constrained and Unconstrained Testing Environment with safe threads for optimization software ([Gould et al., 2015](#)),
 - IMPL © /IMPL-DATA © by Industrial Algorithms Limited, a modeling and solving platform used in the process industries.




Figure 2: Uno’s UML diagram.

Acknowledgments

This work was supported by the Applied Mathematics activity of the U.S. Department of Energy Office of Science, Advanced Scientific Computing Research, under Contract No. DE-AC02-06CH11357. This work was also supported in part by NSF CSSI Grant No. 2104068.

AI usage disclosure

No generative AI was used in the creation of the software, interfaces, implementation of algorithms, or documentation. ChatGPT was used to check spelling, grammar, and clarity of the English text in this paper.

References

- Amestoy, P. R., Duff, I. S., L'Excellent, J.-Y., & Koster, J. (2000). MUMPS: A general purpose distributed memory sparse solver. *International Workshop on Applied Parallel Computing*, 121–130. https://doi.org/10.1007/3-540-70734-4_16
- Andersson, J. A. E., Gillis, J., Horn, G., Rawlings, J. B., & Diehl, M. (2019). CasADi – A software framework for nonlinear optimization and optimal control. *Mathematical Programming Computation*, 11(1), 1–36. <https://doi.org/10.1007/s12532-018-0139-4>
- Baran, M., Bergmann, R., & Przybyś, P. (2026). A Riemannian quasi-Newton algorithm for optimization with Euclidean bounds. *arXiv Preprint arXiv:2605.10573*. <https://doi.org/10.48550/arXiv.2605.10573>
- Bestuzheva, K., Besançon, M., Chen, W.-K., Chmiela, A., Donkiewicz, T., Van Doornmalen, J., Eifler, L., Gaul, O., Gamrath, G., Gleixner, A., & others. (2023). Enabling research through the SCIP optimization suite 8.0. *ACM Transactions on Mathematical Software*, 49(2), 1–21. <https://doi.org/10.1145/3585516>
- Bongartz, I., Conn, A. R., Gould, N., & Toint, P. L. (1995). CUTE: Constrained and unconstrained testing environment. *ACM Transactions on Mathematical Software (TOMS)*, 21(1), 123–160. <https://doi.org/10.1145/200979.201043>
- Büskens, C., & Wassel, D. (2012). The ESA NLP solver WORHP. In *Modeling and optimization in space engineering* (pp. 85–110). Springer. https://doi.org/10.1007/978-1-4614-4469-5_4
- Byrd, R. H., Nocedal, J., & Waltz, R. A. (2006). KNITRO: An integrated package for nonlinear optimization. In *Large-scale nonlinear optimization* (pp. 35–59). Springer. https://doi.org/10.1007/0-387-30065-1_4
- Caillaud, J.-B., Cots, O., Gergaud, J., Martinon, P., & Sed, S. (2024). *OptimalControl.jl: a Julia package to model and solve optimal control problems with ODE's*. <https://doi.org/10.5281/zenodo.13336563>
- Cederberg, D., Zhang, W., Nobel, P., & Boyd, S. (2026). Disciplined nonlinear programming. *arXiv Preprint arXiv:2606.02896*. <https://doi.org/10.48550/arXiv.2606.02896>
- Cesaroni, E., Liuzzi, G., & Lucidi, S. (2026). Derivative-free bilevel optimization with inexact lower-level solutions. *arXiv Preprint arXiv:2603.20759*. <https://doi.org/10.48550/arXiv.2603.20759>
- Desef, B. (2025). Optimization techniques in quantum information. *arXiv Preprint arXiv:2512.15831*. <https://doi.org/10.48550/arXiv.2512.15831>
- Duff, Iain S. (2004). MA57—a code for the solution of sparse symmetric definite and indefinite systems. *ACM Transactions on Mathematical Software (TOMS)*, 30(2), 118–144. <https://doi.org/10.1145/992200.992202>
- Duff, I. S., & Reid, J. K. (1983). The multifrontal solution of indefinite sparse symmetric linear equations. *ACM Trans. Math. Softw.*, 9(3), 302–325. <https://doi.org/10.1145/356044.356047>
- Dunning, I., Huchette, J., & Lubin, M. (2017). JuMP: A modeling language for mathematical

- 214 optimization. *SIAM Review*, 59(2), 295–320. <https://doi.org/10.1137/15M1020575>
- 215 Fletcher, R. (2000). Stable reduced Hessian updates for indefinite quadratic programming.
216 *Mathematical Programming*, 87(2), 251–264. <https://doi.org/10.1007/s101070050113>
- 217 Fletcher, R., & Leyffer, S. (1998). User manual for filterSQP. *Numerical Analysis Report*
218 *NA/181, Department of Mathematics, University of Dundee, Dundee, Scotland*, 35.
- 219 Fourer, R., Gay, D. M., & Kernighan, B. W. (1989). AMPL: A mathematical programming
220 language. In S. W. Wallace (Ed.), *Algorithms and model formulations in mathematical*
221 *programming* (pp. 150–151). Springer. https://doi.org/10.1007/978-3-642-83724-1_12
- 222 Gay, D. M. (1997). *Hooking your solver to AMPL*. Technical Report 93-10, AT&T Bell
223 Laboratories, Murray Hill, NJ, 1993, revised.
- 224 Gill, P. E., Murray, W., & Saunders, M. A. (2005). SNOPT: An SQP algorithm for large-
225 scale constrained optimization. *SIAM Review*, 47(1), 99–131. <https://doi.org/10.1137/S0036144504446096>
- 226
- 227 Gould, N. I., Orban, D., & Toint, P. L. (2015). CUTEst: A constrained and unconstrained
228 testing environment with safe threads for mathematical optimization. *Computational*
229 *Optimization and Applications*, 60(3), 545–557. [https://doi.org/10.1007/s10589-014-](https://doi.org/10.1007/s10589-014-9687-3)
230 [9687-3](https://doi.org/10.1007/s10589-014-9687-3)
- 231 Gruss, L. (2024). *Estimation avancée à horizon glissant, état et paramètres, pour la conduite*
232 *du coeur des réacteurs PWR* [PhD thesis, Ecole nationale supérieure Mines-Télécom
233 Atlantique]. <https://doi.org/10.70675/58a06f96z5f38z4254zad92z861e4b453806>
- 234 Hogg, J. D., Ovtchinnikov, E., & Scott, J. A. (2016). A sparse symmetric indefinite direct
235 solver for GPU architectures. *ACM Transactions on Mathematical Software (TOMS)*,
236 42(1), 1–25. <https://doi.org/10.1145/2756548>
- 237 Hogg, J., & Scott, J. (2010). *An indefinite sparse direct solver for multicore machines*. Tech.
238 Rep. TR-RAL-2010-011, Rutherford Appleton Laboratory, Chilton, Oxfordshire, UK.
- 239 Huangfu, Q., & Hall, J. J. (2018). Parallelizing the dual revised simplex method. *Mathematical*
240 *Programming Computation*, 10(1), 119–142. <https://doi.org/10.1007/s12532-017-0130-5>
- 241 Joshy, A. J., & Hwang, J. T. (2026). modOpt: A modular development environment and
242 library for optimization algorithms. *Advances in Engineering Software*, 213, 104084.
243 <https://doi.org/10.1016/j.advengsoft.2025.104084>
- 244 Kiessling, D., Leyffer, S., & Vanaret, C. (2025). A unified funnel restoration SQP algorithm.
245 *Mathematical Programming*, 1–45. <https://doi.org/10.1007/s10107-025-02284-3>
- 246 Lee, J., & Leyffer, S. (2011). *Mixed integer nonlinear programming*. Springer Science &
247 Business Media. <https://doi.org/10.1007/978-1-4614-1927-3>
- 248 Lewis, F. L., Vrabie, D., & Syrmos, V. L. (2012). *Optimal control*. John Wiley & Sons.
249 <https://doi.org/10.1002/9781118122631>
- 250 Montoisson, A., & Orban, D. (2023). Krylov.jl: A Julia basket of hand-picked Krylov methods.
251 *The Journal of Open Source Software*, 8(89), 5187–5187. [https://doi.org/10.21105/joss.](https://doi.org/10.21105/joss.05187)
252 [05187](https://doi.org/10.21105/joss.05187)
- 253 Nocedal, J., & Wright, S. J. (2006). *Numerical optimization*. Springer. [https://doi.org/10.](https://doi.org/10.1007/978-0-387-40065-5)
254 [1007/978-0-387-40065-5](https://doi.org/10.1007/978-0-387-40065-5)
- 255 Shin, S., Anitescu, M., & Pacaud, F. (2024). Accelerating optimal power flow with GPUs: SIMD
256 abstraction of nonlinear programs and condensed-space interior-point methods. *Electric*
257 *Power Systems Research*, 236, 110651. <https://doi.org/10.1016/j.epsr.2024.110651>
- 258 Sra, S., Nowozin, S., & Wright, S. J. (2011). *Optimization for machine learning*. MIT press.
259 <https://doi.org/10.7551/mitpress/8996.001.0001>

- 260 Vanaret, C., & Leyffer, S. (2026). Implementing a unified solver for nonlinearly constrained
261 optimization. *Mathematical Programming Computation*. [https://doi.org/10.1007/s12532-](https://doi.org/10.1007/s12532-026-00310-9)
262 [026-00310-9](https://doi.org/10.1007/s12532-026-00310-9)
- 263 Wächter, A., & Biegler, L. T. (2006). On the implementation of an interior-point filter
264 line-search algorithm for large-scale nonlinear programming. *Mathematical Programming*,
265 *106*(1), 25–57. <https://doi.org/10.1007/s10107-004-0559-y>
- 266 Wu, E., Kenway, G., Mader, C. A., Jasa, J., & Martins, J. R. R. A. (2020). pyOptSparse:
267 A Python framework for large-scale constrained nonlinear optimization of sparse systems.
268 *Journal of Open Source Software*, *5*(54), 2564. <https://doi.org/10.21105/joss.02564>

DRAFT